
PROOF OVER PROSE: TRAINING LANGUAGE MODELS ON FORMAL MULTI-HOP DERIVATIONS

Anonymous authors

Paper under double-blind review

ABSTRACT

Language models learn to reason from examples of reasoning, and those examples are almost always written in prose. We ask whether the notation matters. Our data comes from a generator that samples multi-hop deduction problems over Horn clauses, where every step offers four locally plausible continuations, exactly one of which follows from the current state, and the answer is unique by construction. Each latent proof is written out twice. One version is a formal derivation with numbered lines and named inference rules. The other is a controlled English explanation of the same steps. Prompt, answer and proof graph are identical within a pair, so two models trained on opposite sides differ only in the notation they were shown.

Which notation produces the better reasoner depends on how many hops the training proofs contain. Through five to twenty hops the English rendering generalizes further to unseen depths. At twenty-five hops the ordering reverses and formal supervision reaches 75.0 ± 9.5 percent out-of-distribution accuracy where English reaches 17.8 ± 4.7 . Drawing sixteen samples instead of one moves the reversal ten hops earlier, to fifteen, so formal supervision first widens the set of problems a model can solve at all and only later improves its first attempt. A formal derivation can be checked mechanically, so the wider set is usable without an answer key, and returning the first sample a proof checker accepts yields 97.3 ± 1.1 percent accuracy where the same procedure on English traces yields 33.8 ± 1.4 percent.

The effect is not confined to the controlled setting. We midtrain Qwen2.5-7B on 2.5 billion tokens with ten percent of the mixture replaced by these derivations, and accuracy improves on held-out ProofWriter problems that the generator never produced. Replacing the same ten percent with ordinary long documents, matched to the formal traces in length distribution, leaves accuracy at the control level, so document length alone does not explain the gain.

1 INTRODUCTION

A language model that answers a multi-hop question correctly has, somewhere in its training data, seen many worked examples of chains of inference. Those examples are written in prose, because that is how people write when they explain themselves. The same derivation could instead be written in formal logic, with each line naming the rule that produced it and the earlier lines it depends on. Whether that changes what the model learns is an empirical question that the literature has not asked, since the two notations never appear over the same underlying problems in natural corpora. Two solutions written in different styles are rarely two renderings of one computation, so any comparison drawn from found data confounds notation with problem difficulty and with the care of whoever wrote the solution.

We remove that confound by generating both renderings from a single source. A deterministic generator samples a latent proof over Horn clauses and then writes it out twice, once as a formal derivation and once as a controlled English explanation. The prompt, the answer, the premises and the proof graph are shared. Only the notation of the supervised trace changes, so a difference between two models trained on opposite sides of a pair can be attributed to the notation (Figure 1). No prover runs at inference time, and both conditions generate under identical token budgets, so we are measuring what the notation teaches and any difference cannot come from access to a solver.




Figure 1: The paired-trace construction. A generator samples a latent proof and renders it as a formal derivation or as a controlled natural-language explanation. The prompt, the answer and the proof graph are shared, and only the supervised trace differs. Models are trained on one rendering and evaluated beyond the training depth range.

Our claim is not that a model trained on formal derivations begins reasoning in formal logic internally. What the formal rendering supplies is an explicit, uniform record of how a multi-hop derivation is assembled, in which every step states which premises it consumed and which rule licensed it. A model trained on prose sees the same chain with those dependencies left implicit in the grammar of the sentences. The question is whether making the structure explicit produces a better multi-hop reasoner, and the answer turns out to depend on how long the chains in training were.

Through training ranges of five to twenty hops the English rendering is reliably stronger on held-out problems deeper than anything seen in training. At twenty-five hops the comparison reverses by a factor of four, with formal supervision at 75.0 ± 9.5 percent against 17.8 ± 4.7 (Table 1). Neither target length nor the number of supervised tokens accounts for the reversal, and an independently generated constraint-propagation task reproduces it.

How the model is decoded moves the reversal. Greedy decoding places it at twenty-five hops. Sixteen samples place it at fifteen, where formal supervision reaches 69.9 percent against 56.9. In the range between those two points a formally trained model can already construct a correct derivation within a modest sample budget while still failing on its first attempt. Any study that reported only greedy accuracy would locate the effect ten hops too late and would describe the intervening range as showing no formal advantage.

That gap between what a model can produce and what it produces first is where formal notation pays. Validity of a formal derivation is decidable by the same engine that generated the problem, so a sample can be checked without knowing the answer. Nearly every test prompt admits a correct and mechanically valid proof within sixteen samples, at 98.1 ± 0.5 percent, and returning the first sample the checker accepts recovers 97.3 ± 1.1 percent single-answer accuracy. Applied to English traces the identical procedure recovers 33.8 ± 1.4 percent, largely because long English derivations often fail to terminate before the budget runs out, leaving nothing for the checker to accept. Recent work argues that reinforcement learning with verifiable rewards mostly reweights a model toward completions it can already sample (Yue et al., 2025). If that is right, then the breadth we measure here is the quantity that bounds what such posttraining can reach.

Controlled fine-tuning on synthetic problems cannot tell us whether any of this matters for a language model trained the ordinary way. We therefore midtrain Qwen2.5-7B on 2.5 billion tokens of a production mixture, replacing ten percent of the loss tokens with matched derivations at twenty-five hops and holding every other hyperparameter fixed across conditions. Both renderings improve accuracy on held-out ProofWriter problems. A corpus of ordinary long documents matched to the formal traces in length distribution does not, which rules out the most obvious alternative explanation. The sampled and greedy orderings separate here as they did in the controlled study.

2 RELATED WORK

Three lines of work bear on the question we ask, and none of them varies the notation of a reasoning trace while holding the reasoning fixed.

The first established that intermediate steps help. Prompting a model to write out its reasoning improves multi-step accuracy (Wei et al., 2022; Kojima et al., 2022), and the idea has been extended by bootstrapping rationales (Zelikman et al., 2022), decomposing questions (Zhou et al., 2023), searching over partial reasoning states (Yao et al., 2023) and supervising a scratchpad directly (Nye et al., 2021). Theory explains part of the gain, since intermediate tokens let a constant-depth transformer carry out serial computation that it cannot perform in a single forward pass (Feng et al., 2023; Merrill & Sabharwal, 2024; Li et al., 2024). What this work holds constant is the notation. The steps are written in prose throughout, which leaves open whether prose is the reason the steps help or merely the medium they arrived in.

The second line asks what is actually being learned, and it reaches for synthetic data to do so. Controlling the generator makes it possible to ask questions that benchmark accuracy cannot answer, as in the iGSM family (Ye et al., 2025a;b) and in work that controls the generating grammar outright (Allen-Zhu & Li, 2023). The findings are sobering. Models trained to reason can satisfy the objective through statistical shortcuts (Zhang et al., 2023), they fail on compositions that scale does not fix (Dziri et al., 2023), and formal analysis of their derivations exposes error modes that aggregate scores conceal (Saparov & He, 2023; Clark et al., 2020; Tafjord et al., 2021). The locality of the training distribution has even been proposed as the reason intermediate steps help at all (Prystawski et al., 2023). These studies vary the problems, the model and the training signal. The representation of the trace itself remains a free variable that nobody has moved.

The third line treats formality as an inference-time resource. Program-aided prompting delegates computation to an interpreter (Gao et al., 2023; Chen et al., 2023), logic pipelines hand the problem to a solver (Pan et al., 2023; Ling et al., 2023), theorem provers construct objects a kernel will check (Polu & Sutskever, 2020; Jiang et al., 2023; Yang et al., 2023; Song et al., 2024), autoformalization translates informal statements into such systems (Wu et al., 2022), and recent proposals type a trace to make it verifiable (Perrier, 2025). In all of these the formal object earns its place at test time, through an external checker or an external solver. Whether formality is worth anything as *training* signal, with no solver present at inference, is a different question, and it is the one we take up.

Our measurements also touch a live disagreement about sampling. Repeated sampling converts inference compute into solved problems (Brown et al., 2024), verifiers and process supervision select among those samples (Cobbe et al., 2021; Lightman et al., 2024), self-consistency selects by agreement (Wang et al., 2023), and how best to spend a fixed budget has been studied both empirically (Snell et al., 2025; Muennighoff et al., 2025) and theoretically (Setlur et al., 2025). Against that, one recent argument holds that verifiable-reward training mostly sharpens a model toward what it could already sample (Yue et al., 2025; Guo et al., 2025; Shao et al., 2024; Gandhi et al., 2025). Our selector needs no trained verifier and no answer key, since validity is decidable by the generating engine, and the breadth it exploits is precisely the quantity that argument says bounds posttraining.

Two further concerns frame how we report. Prose rationales can omit what actually drove a prediction (Turpin et al., 2023; Lanham et al., 2023), so we separate answer accuracy from proof validity everywhere. And systematic generalization is conventionally probed by controlling the train and test distributions (Lake & Baroni, 2018; Keysers et al., 2020; Hupkes et al., 2020; Anil et al., 2022), which is what our depth split does, with a spurious-marker variant to test shortcut reliance (Geirhos et al., 2020).

3 DATA AND EXPERIMENTAL SETUP

Each example begins with a latent directed acyclic proof. A generator emits a prompt containing ground facts, rules, a query and distractor branches, then traverses the supporting path to produce the answer. A rendering function writes that path in one of two trace languages. The FORMAL rendering declares constants and predicates and applies named inference rules to numbered lines. The NATURAL rendering expresses the same substitutions and conclusions, line for line, in controlled




Figure 2: Out-of-distribution answer accuracy on BRANCHPROOF against maximum training depth, for three decoding rules. Shaded regions give one standard deviation over three seeds. Greedy decoding and pass@1 place the crossover at maximum training depth 25. Sampling with sixteen draws places it at 15, ten depth levels earlier.

sentences. Both targets end with the same answer field and the prompt is natural language in both conditions.

BRANCHPROOF composes Horn-style inferences in which every step offers four locally plausible branches, exactly one derivable from the current state, with a fresh constant introduced at every layer. The generator computes the full Horn closure of each example and accepts it only when exactly one candidate answer is derivable, which removes any reward for guessing among defensible answers. ATTRCON propagates values among object attributes under generated constraints, and its vocabulary and proof construction are disjoint from BRANCHPROOF. Generation parameters are held fixed across every controlled experiment, at branching factor 4 and distractor ratio 0.5, with depths sampled round-robin. Corpus seeds are disjoint across training, evaluation and sampling roles. Appendices G to F give paired targets, the grammar, the generator and the validators, and Appendix A reports the answer-uniqueness validation.

The controlled experiment fine-tunes OLMo-3-7B (OLMo Team et al., 2025) with LoRA on 50,000 materialized examples per condition, with maximum training depth in $\{5, 10, 15, 20, 25\}$ and three seeds per condition. Evaluation covers depths $\{1, 2, 5, 10, 12, 15, 18, 20, 25, 30, 35, 40, 45, 50\}$ with 32 prompts per depth, and at each prompt sixteen samples at temperature 1.0 together with one greedy decode. Sampled coverage uses the unbiased pass@k estimator (Chen et al., 2021). Both trace languages receive an identical 7,168-token generation limit inside a 16,384-token window, and a tokenizer audit confirms every gold prompt and target fits in both renderings, so neither condition is truncated. Depths above the maximum training depth are out of distribution. Joint metrics require the correct answer together with a mechanically valid proof, where formal traces are checked directly by the proof engine and natural traces by a deterministic translation into the same engine, so answer-only accuracy is reported alongside every joint metric. Appendices D and H give the remaining configuration and the metric definitions.

4 HOW PROOF DEPTH CHANGES THE COMPARISON

Out-of-distribution answer accuracy as a function of maximum training depth is given in Table 1. Through maximum training depth 20 the natural condition is ahead on every decoding rule. At maximum training depth 25 the formal condition reaches 0.750 pass@1 against 0.178, a factor of four, and 0.996 pass@16 against 0.425.

The two decoding rules do not place the reversal at the same depth. Greedy accuracy and pass@1 both cross at maximum training depth 25, while pass@16 crosses at maximum training depth 15, where the formal condition reaches 0.699 against 0.569. The interval between the two crossings is a regime in which formal supervision has already produced a model that can construct a correct derivation within sixteen attempts, and has not yet produced one that does so on the first attempt.

Table 1: Out-of-distribution answer accuracy on BRANCHPROOF by maximum training depth, three seeds per cell. Greedy and pass@1 order the conditions identically. pass@16 reverses ten depth levels earlier.

metric and condition	maximum training depth				
	5	10	15	20	25
greedy, FORMAL	0.108	0.331	0.281	0.220	0.821
greedy, NATURAL	0.563	0.600	0.574	0.684	0.487
pass@1, FORMAL	0.098	0.308	0.261	0.232	0.750
pass@1, NATURAL	0.524	0.580	0.495	0.312	0.178
pass@16, FORMAL	0.500	0.678	0.699	0.804	0.996
pass@16, NATURAL	0.735	0.711	0.569	0.453	0.425

Coverage is acquired before single-sample reliability. A study reporting greedy accuracy alone would locate the effect ten depth levels too high and would describe the intermediate range as showing no formal advantage at all.

The natural condition degrades monotonically in pass@16 from 0.735 at depth 5 to 0.425 at depth 25 while the formal condition rises monotonically from 0.500 to 0.996, so the two renderings respond to increasing depth coverage in opposite directions.

Length and budget controls. Formal and natural targets differ in length, which raises the possibility that the comparison measures supervision quantity. Two controls address this. Token budgets at generation are identical by construction, and the tokenizer audit confirms no truncation in either condition. Matched-token training, in which the number of supervised target tokens is equated across conditions, leaves the ordering unchanged. Details are in Appendix I.

Replication on an independent task. ATTRCON shares no vocabulary and no proof construction with BRANCHPROOF. The correctness advantage of formal supervision at high training depth reproduces there, which makes it unlikely that the effect depends on a particular rule system.

5 CHECKING SAMPLES WITHOUT AN ANSWER KEY

A model whose sixteen samples contain a correct proof is only useful if that proof can be picked out without already knowing the answer. For formal derivations it can be, since validity is decidable by the same engine that generated the problem. Consider a selector that draws k samples, checks each in order and returns the answer of the first one the checker accepts.

Proposition 1. *Let each sample be, independently, valid and correct with probability a , valid and incorrect with probability b , and rejected with probability $1 - a - b$. The first-valid selector returns a correct answer with probability $(1 - (1 - a - b)^k) a / (a + b)$. If the checker admits no incorrect proof, so that $b = 0$, this equals $1 - (1 - a)^k$, the joint pass@ k .*

The proof is elementary and appears in Appendix L. Two measurable quantities therefore bound selection, joint pass@ k from above and checker soundness $a / (a + b)$ independently of k . On BRANCHPROOF at maximum training depth 25 the formal condition attains joint pass@16 of 98.1 ± 0.5 percent and the selector attains 97.3 ± 1.1 percent, so the 0.8 point gap is the empirical cost of $b > 0$ and the selector sits close to the permitted ceiling.

The same selector applied to translated natural-language traces recovers 33.8 ± 1.4 percent. The deficit is not primarily a failure of translation. Long natural traces frequently fail to terminate inside the generation budget, so no checkable derivation exists to accept, which drives $a + b$ down and leaves the selector with nothing to select from. No reward model and no gold labels enter the procedure.

6 WHERE THE EFFECT REVERSES

Capacity reorders the comparison, since the two conditions are close on Qwen2.5-7B at 69.9 ± 11.1 against 74.9 ± 7.0 pass@1 and the natural condition leads on OLMo-3-32B at 98.9 ± 0.8 against 63.9 ± 5.1 , where formal coverage nonetheless stays near ceiling at 98.8 ± 0.9 pass@16. Concatenating both renderings into one target fails outright, below one percent, while mode-conditioned training at 32B lifts the formal mode to 85.0 ± 8.7 , which remains below single-modality prose at that scale. A spurious schema marker planted on 80 percent of training examples and removed at test time makes the natural condition the robust one, at 78.5 ± 14.9 against 43.1 ± 7.7 , so the formal advantage is not generic robustness to distribution shift. Models trained from scratch at 50M to 200M parameters fail in both conditions at every size, so the effect operates on competence supplied by pretraining and does not create it. Full tables for all four experiments are in Appendix J.

7 MIDTRAINING A 7B MODEL ON FORMAL DERIVATIONS

Whether the same variable matters when derivations are a small minority of an ordinary training mixture is tested by midtraining Qwen2.5-7B (Yang et al., 2024) on a production corpus with a fixed fraction replaced by matched derivations, on a different model family and against benchmarks the generator never saw.

7.1 DESIGN

Five conditions share every hyperparameter, every data path outside the replaced slice, and an identical token budget of 2,500,853,760 training tokens, reached as 2,385 steps of global batch 128 at sequence length 8,192. The conditions differ only in the ten percent of loss tokens that is replaced.

1. **Control** replaces nothing and trains on the unmodified mixture.
2. **LongDoc** replaces ten percent with long documents drawn from the same production corpus, selected to match the token-length histogram of the formal traces and carrying no deep deduction.
3. **Logic** replaces ten percent with formal derivations at depth 25.
4. **Natural** replaces ten percent with the natural-language rendering of the same latent proofs.
5. **Condensed** replaces ten percent with a compact formal rendering of the same latent proofs.

LongDoc separates two explanations of any observed gain, since formal derivations are long and long documents change what a model learns about maintaining state across a context. Matching the length histogram while removing the deduction isolates the second factor. Median length is 3,712 tokens against 3,750 for the formal traces, mean 3,832 against 3,856, maximum 7,750 against 7,728, and all thirty populated 250-token bins agree within about three percent.

7.2 CORPORA AND PACKING

A corpus of 72,000 examples is generated with depths 1 to 25 round-robin, branching factor 4 and distractor ratio 0.5, matching the controlled study. The three trace corpora render the same latent proofs, so the comparison between them holds the underlying computation fixed as in Section 3. Rendered lengths and packing statistics appear in Appendix K.

A document longer than one window is excluded and counted, and none is split across windows, so the mid-derivation truncation that would destroy the depth structure under study cannot occur. Zero documents were excluded and zero split in every condition. Padding uses a dedicated token masked out of the loss. A decoded-batch audit verifies both properties on real loader windows and training refuses to start if it fails.

Because the replaced slice is fixed in tokens and the condensed rendering is 2.6 times shorter, the condensed condition traverses its proof set 2.31 times where Logic and Natural traverse theirs 0.90 and 0.89 times. Any comparison involving the condensed condition therefore confounds rendering with exposure, and we treat it accordingly in Section 7.4.

Table 2: Exact match on the external ProofWriter items by inference depth, 500 items per depth. LongDoc is length-matched to the formal traces and carries no deep deduction. Depth 0 requires lookup and no inference.

inference depth	0	1	2	3	5
Control	0.442	0.328	0.440	0.444	0.460
LongDoc	0.458	0.322	0.444	0.454	0.440
Logic	0.480	0.358	0.494	0.500	0.502
Natural	0.488	0.388	0.562	0.570	0.564
Condensed	0.460	0.346	0.460	0.480	0.474

Table 3: Formal minus natural on the near-transfer derivation family, by inference depth and decoding rule. Positive favours the formal condition. The sign of the contrast follows the decoding rule, with no clear depth trend.

decoding	d5	d10	d15	d20	d25
greedy	-0.175	-0.165	-0.055	-0.040	-0.105
pass@16	+0.035	+0.070	+0.140	+0.040	-0.035
maj@16	-0.140	-0.030	-0.005	+0.025	-0.070

7.3 INSTRUCTION TUNING AND EVALUATION

All five checkpoints receive an identical instruction-tuning stage of one epoch over 100,000 examples at learning rate 5×10^{-6} , full-parameter under FSDP. Evaluation uses two families. The external family is ProofWriter (Tafjord et al., 2021) under the open-world assumption at inference depths 0, 1, 2, 3 and 5 with 500 items per depth, which the generator never saw. The near-transfer family asks for a derivation on held-out problems from the same generator at depths 5, 10, 15, 20 and 25 with 200 items per depth. The metric is exact match on the extracted answer. Sampled measurements draw sixteen completions at temperature 0.8 and top- p 0.95, replaying the greedy prompts so the two decodings differ in the sampler alone. Appendix K gives every remaining setting.

7.4 RESULTS

On the external ProofWriter items both trace conditions improve over Control and over LongDoc, and LongDoc does not improve over Control (Table 2). At inference depth 0, which requires lookup and no inference, all five conditions fall between 0.442 and 0.488, and the separation appears only once inference is required. That is the shape a deduction effect should have, and it is not reproduced by exposure to long documents alone.

The near-transfer family separates the conditions far more sharply. Control sits near the floor at 0.000 to 0.040 exact match across depths, where every condition trained on derivations reaches 0.23 to 0.29, so the replaced slice teaches a capability the control mixture does not supply at all.

Under greedy decoding the natural condition is ahead of the formal condition on thirteen of fifteen tasks. Under sampling the ordering reverses. Table 3 gives the contrast under both decodings.

Pooling over depths and bootstrapping over documents with 4,000 resamples gives the contrasts in Table 4. Both trace conditions and LongDoc separate from Control, the trace conditions separate from LongDoc, and the formal condition exceeds the natural condition by 0.050 in pass@16 and by 0.025 in mean sampled accuracy.

Formal supervision is behind under greedy decoding and ahead under sampling, the same pattern that separated the two crossover depths in Section 4, now on a different model family with the derivations forming a tenth of an ordinary mixture. The proportion of greedy generations reaching a parseable answer is 0.645 to 0.785 for Control and 0.940 to 1.000 for the trace conditions, so the conditions differ in whether a derivation is completed at all.

Table 4: Paired bootstrap over documents, 4,000 resamples, pass@16 pooled over inference depths 5 to 25. Intervals cover evaluation sampling error only. Section 7.5 treats training-run variance separately.

contrast	difference	95 percent interval
Condensed – Control	+0.276	[0.238, 0.313]
Logic – Control	+0.226	[0.186, 0.265]
Natural – Control	+0.176	[0.139, 0.214]
LongDoc – Control	+0.138	[0.104, 0.171]
Logic – LongDoc	+0.088	[0.050, 0.125]
Condensed – Natural	+0.100	[0.064, 0.136]
Logic – Natural	+0.050	[0.013, 0.087]

The condensed condition attains the largest gain over Control and exceeds the formal condition by 0.050. Its documents are 2.6 times shorter, so its advantage is not a length effect. It traverses its proof set 2.6 times more often, so its advantage is confounded with exposure and cannot be attributed to the rendering. Resolving that requires an exposure-matched condensed corpus, which is not part of this work.

7.5 HOW FAR THIS RESULT GOES

Every midtraining condition was trained once, so the intervals in Table 4 estimate evaluation sampling error. They say nothing about the variability that retraining under a different seed would produce. That second quantity is measurable. An auxiliary experiment on the same evaluation suite, described in Appendix B, gives a per-run standard deviation near 0.019 and therefore near 0.027 for a difference between two independently seeded conditions. Combined with the evaluation term this gives about 0.033 for the formal against natural contrast, whose interval then becomes approximately $[-0.015, +0.115]$ and includes zero.

The contrasts against Control, between 0.176 and 0.276, are six to ten times the training standard deviation and are not at risk from single-seed training. The contrast between the two trace languages, at 0.050, is roughly twice that standard deviation and is a sign reversal against the greedy ordering, which is the pattern most likely to be produced by run-to-run variation. We therefore report the transfer-scale formal advantage as consistent with the controlled study. We do not claim it is independently established. Replication under a second instruction-tuning seed is the appropriate remedy, since enlarging the evaluation set would shrink the smaller error term and leave the larger untouched.

The depth of the replaced derivations was fixed at 25 in every condition, so the transfer study locates no threshold of its own and inherits the one estimated from OLMo-3-7B in Section 4. An earlier configuration at sequence length 4,096, with a five percent replacement share and derivations capped at depth 14, produced no measurable transfer. It differs from the present configuration in window, replacement share, rendering and evaluation suite, so it motivates this design and is not a controlled below-threshold comparison.

8 DISCUSSION

Neither notation is better in general. Which one produces the stronger reasoner flips as the training proofs get deeper, and where it flips depends on how the model is decoded. A study reporting greedy accuracy alone places the crossover at depth 25 and concludes that formal supervision is harmful everywhere below it, where the same experiment read at pass@16 places it at depth 15 and shows the formal condition ahead through the upper half of the range. The two readings measure different things, and the second is the one relevant to any system drawing more than one sample.

That distinction connects these measurements to verifiable-reward posttraining. If such training principally reweights toward completions the base model can already produce (Yue et al., 2025), then coverage after midtraining bounds what the later stage reaches. Trace language moves that coverage by between 0.050 and 0.276 pass@16 in our transfer setting, placing it upstream of the posttraining

stage. Verifying the ordering requires running that stage on both trace conditions, which we have not done.

9 LIMITATIONS

The controlled study uses one model family at 7B and one at 32B, and the transfer study uses a third. Depth thresholds estimated on OLMo-3-7B are applied to Qwen2.5-7B by assumption. We did not measure them there.

The transfer conditions were trained once each, with the consequences set out in Section 7.5.

Proof validity for natural-language traces is established through a deterministic translation into the proof engine. The translation is a component that the formal condition does not require, and although answer-only accuracy is reported alongside every joint metric, the joint natural-language numbers carry that dependency.

The generator produces one family of Horn-style deduction problems and one constraint-propagation family. Neither involves arithmetic, natural-language ambiguity or retrieval, and the depth axis studied here is proof depth. Context length and problem breadth are untouched.

The condensed rendering is confounded with exposure, as described in Section 7.4.

10 CONCLUSION

Holding a derivation fixed and changing only the language in which it is written changes what a model learns from it, and the direction is set by the depth of the derivations in the training range. Formal supervision acquires coverage before single-sample accuracy, placing two thresholds ten depth levels apart and making the decoding rule part of the experimental design. Because formal derivations are checkable, that coverage converts into single-answer accuracy without labels at a rate Proposition 1 bounds and measurement approaches. The same dissociation appears when matched derivations form a tenth of an ordinary midtraining mixture, against both an untreated control and a length-matched control without deduction.

REFERENCES

- Zeyuan Allen-Zhu and Yuanzhi Li. Physics of language models: Part 1, learning hierarchical language structures. *arXiv preprint arXiv:2305.13673*, 2023. doi: 10.48550/arXiv.2305.13673.
- Cem Anil, Yuhuai Wu, Anders Andreassen, Aitor Lewkowycz, Vedant Misra, Vinay Ramasesh, Ambrose Slone, Guy Gur-Ari, Ethan Dyer, and Behnam Neyshabur. Exploring length generalization in large language models. In *Advances in Neural Information Processing Systems 35*, 2022.
- Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*, 2024. doi: 10.48550/arXiv.2407.21787.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021. doi: 10.48550/arXiv.2107.03374.
- Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen. Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks. *Transactions on Machine Learning Research*, 2023.
- Peter Clark, Oyvind Tafjord, and Kyle Richardson. Transformers as soft reasoners over language. In *Proceedings of the 29th International Joint Conference on Artificial Intelligence*, 2020.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. doi: 10.48550/arXiv.2110.14168.

486 Nouha Dziri, Ximing Lu, Melanie Sclar, Xiang Lorraine Li, Liwei Jiang, Bill Yuchen Lin, Sean
487 Welleck, Peter West, Chandra Bhagavatula, Ronan Le Bras, et al. Faith and fate: Limits of
488 transformers on compositionality. In *Advances in Neural Information Processing Systems* 36,
489 2023.

490 Guhao Feng, Bohang Zhang, Yuntian Gu, Haotian Ye, Di He, and Liwei Wang. Towards revealing
491 the mystery behind chain of thought: A theoretical perspective. In *Advances in Neural Information*
492 *Processing Systems* 36, 2023.

493 Kanishk Gandhi, Ayush Chakravarthy, Anikait Singh, Nathan Lile, and Noah D. Goodman. Cognitive
494 behaviors that enable self-improving reasoners, or, four habits of highly effective STaRs. In
495 *Conference on Language Modeling*, 2025. doi: 10.48550/arXiv.2503.01307.

496 Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and
497 Graham Neubig. PAL: Program-aided language models. In *Proceedings of the 40th International*
498 *Conference on Machine Learning*, 2023.

500 Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias
501 Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine*
502 *Intelligence*, 2(11):665–673, 2020. doi: 10.1038/s42256-020-00257-z.

504 Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu,
505 Shirong Ma, Peiyi Wang, Xiao Bi, et al. DeepSeek-R1 incentivizes reasoning in LLMs through
506 reinforcement learning. *Nature*, 645(8081):633–638, 2025. doi: 10.1038/s41586-025-09422-z.

507 Dieuwke Hupkes, Verna Dankers, Mathijs Mul, and Elia Bruni. Compositionality decomposed: How
508 do neural networks generalise? *Journal of Artificial Intelligence Research*, 67:757–795, 2020. doi:
509 10.1613/jair.1.11674.

511 Albert Q. Jiang, Sean Welleck, Jin Peng Zhou, Wenda Li, Jiacheng Liu, Mateja Jamnik, Timothée
512 Lacroix, Yuhuai Wu, and Guillaume Lample. Draft, sketch, and prove: Guiding formal theorem
513 provers with informal proofs. In *International Conference on Learning Representations*, 2023. doi:
514 10.48550/arXiv.2210.12283.

515 Daniel Keysers, Nathanael Schärli, Nathan Scales, Hylke Buisman, Daniel Furrer, Sergii Kashubin,
516 Nikola Momchev, Danila Sinopalnikov, Lukasz Stafiniak, Tibor Tihon, Dmitry Tsarkov, Xiao Wang,
517 Marc van Zee, and Olivier Bousquet. Measuring compositional generalization: A comprehensive
518 method on realistic data. In *International Conference on Learning Representations*, 2020.

519 Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large
520 language models are zero-shot reasoners. In *Advances in Neural Information Processing Systems*
521 35, 2022. doi: 10.48550/arXiv.2205.11916.

522 Brenden M. Lake and Marco Baroni. Generalization without systematicity: On the compositional
523 skills of sequence-to-sequence recurrent networks. In *Proceedings of the 35th International*
524 *Conference on Machine Learning*, 2018.

526 Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny
527 Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, et al. Measuring
528 faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023. doi:
529 10.48550/arXiv.2307.13702.

530 Zhiyuan Li, Hong Liu, Denny Zhou, and Tengyu Ma. Chain of thought empowers transformers to
531 solve inherently serial problems. In *International Conference on Learning Representations*, 2024.
532 doi: 10.48550/arXiv.2402.12875.

533 Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan
534 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International*
535 *Conference on Learning Representations*, 2024. doi: 10.48550/arXiv.2305.20050.

536 Zhan Ling, Yunhao Fang, Xuanlin Li, Zhiao Huang, Mingu Lee, Roland Memisevic, and Hao
537 Su. Deductive verification of chain-of-thought reasoning. In *Advances in Neural Information*
538 *Processing Systems* 36, 2023.

-
- William Merrill and Ashish Sabharwal. The expressive power of transformers with chain of thought. In *International Conference on Learning Representations*, 2024.
- Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori Hashimoto. s1: Simple test-time scaling. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025. doi: 10.48550/arXiv.2501.19393.
- Maxwell Nye, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, Jacob Austin, David Bieber, David Dohan, Aitor Lewkowycz, Maarten Bosma, David Luan, Charles Sutton, and Augustus Odena. Show your work: Scratchpads for intermediate computation with language models. *arXiv preprint arXiv:2112.00114*, 2021. doi: 10.48550/arXiv.2112.00114.
- OLMo Team, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, Akshita Bhagia, Yuling Gu, Shengyi Huang, Matt Jordan, et al. 2 OLMo 2 furious. In *Conference on Language Modeling*, 2025. doi: 10.48550/arXiv.2501.00656.
- Liangming Pan, Alon Albalak, Xinyi Wang, and William Yang Wang. Logic-LM: Empowering large language models with symbolic solvers for faithful logical reasoning. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. doi: 10.18653/v1/2023.findings-emnlp.248.
- Elija Perrier. Typed chain-of-thought: A curry-howard framework for verifying LLM reasoning. *arXiv preprint arXiv:2510.01069*, 2025. doi: 10.48550/arXiv.2510.01069.
- Stanislas Polu and Ilya Sutskever. Generative language modeling for automated theorem proving. *arXiv preprint arXiv:2009.03393*, 2020. doi: 10.48550/arXiv.2009.03393.
- Ben Prystawski, Michael Y. Li, and Noah D. Goodman. Why think step by step? reasoning emerges from the locality of experience. In *Advances in Neural Information Processing Systems 36*, 2023.
- Abulhair Saparov and He He. Language models are greedy reasoners: A systematic formal analysis of chain-of-thought. In *International Conference on Learning Representations*, 2023.
- Amrith Setlur, Nived Rajaraman, Sergey Levine, and Aviral Kumar. Scaling test-time compute without verification or RL is suboptimal. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025. doi: 10.48550/arXiv.2502.12118.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. doi: 10.48550/arXiv.2402.03300.
- Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM test-time compute optimally can be more effective than scaling parameters for reasoning. In *International Conference on Learning Representations*, 2025. doi: 10.48550/arXiv.2408.03314.
- Peiyang Song, Kaiyu Yang, and Anima Anandkumar. Lean copilot: Large language models as copilots for theorem proving in lean. *arXiv preprint arXiv:2404.12534*, 2024. doi: 10.48550/arXiv.2404.12534.
- Oyvind Tafjord, Bhavana Dalvi, and Peter Clark. ProofWriter: Generating implications, proofs, and abductive statements over natural language. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, 2021. doi: 10.18653/v1/2021.findings-acl.317.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems 36*, 2023. doi: 10.48550/arXiv.2305.04388.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations*, 2023. doi: 10.48550/arXiv.2203.11171.

-
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems 35*, 2022. doi: 10.48550/arXiv.2201.11903.
- Yuhuai Wu, Albert Q. Jiang, Wenda Li, Markus N. Rabe, Charles Staats, Mateja Jamnik, and Christian Szegedy. Autoformalization with large language models. In *Advances in Neural Information Processing Systems 35*, 2022.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. doi: 10.48550/arXiv.2412.15115.
- Kaiyu Yang, Aidan Swope, Alex Gu, Rahul Chalamala, Peiyang Song, Shixing Yu, Saad Godil, Ryan Prenger, and Anima Anandkumar. LeanDojo: Theorem proving with retrieval-augmented language models. In *Advances in Neural Information Processing Systems 36 Datasets and Benchmarks Track*, 2023.
- Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. In *Advances in Neural Information Processing Systems 36*, 2023.
- Tian Ye, Zicheng Xu, Yuanzhi Li, and Zeyuan Allen-Zhu. Physics of language models: Part 2.1, grade-school math and the hidden reasoning process. In *International Conference on Learning Representations*, 2025a. doi: 10.48550/arXiv.2407.20311.
- Tian Ye, Zicheng Xu, Yuanzhi Li, and Zeyuan Allen-Zhu. Physics of language models: Part 2.2, how to learn from mistakes on grade-school math problems. In *International Conference on Learning Representations*, 2025b. doi: 10.48550/arXiv.2408.16293.
- Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang. Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model? In *Advances in Neural Information Processing Systems 39*, 2025. doi: 10.48550/arXiv.2504.13837.
- Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. STaR: Bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems 35*, 2022.
- Honghua Zhang, Liunian Harold Li, Tao Meng, Kai-Wei Chang, and Guy Van den Broeck. On the paradox of learning to reason from data. In *Proceedings of the 32nd International Joint Conference on Artificial Intelligence*, 2023.
- Denny Zhou, Nathanael Schärli, Le Hou, Jason Wei, Nathan Scales, Xuezhi Wang, Dale Schuurmans, Claire Cui, Olivier Bousquet, Quoc V. Le, and Ed H. Chi. Least-to-most prompting enables complex reasoning in large language models. In *International Conference on Learning Representations*, 2023.

A ANSWER-UNIQUENESS VALIDATION OF THE GENERATOR

Every metric in this paper compares a generated answer against a single labelled answer, which is only meaningful if the labelled answer is the sole one derivable from the premises. The generator enforces this by construction and the property is verified independently.

Each example introduces fresh constants c_0 through c_{depth} , so no constant is reused across layers and no two branches can derive the same atom through coincidental reuse of a state predicate. A preflight check computes the full Horn closure of 1,000 training instances and 2,000 evaluation instances with a solver independent of the generator, and records the number of candidate final states derivable in each. The unique-answer rate is 1.0, the maximum derived-candidate count is 1, there are no generation failures, and answer positions are balanced across the four candidates.

Token headroom is verified over the same instances. Evaluation uses a 16,384-token context with a 7,168-new-token generation cap applied identically to both trace languages. Across every validation row and every depth, the minimum measured headroom is 494 tokens for generation and 2,788 tokens for the full context, so neither trace language is truncated at any depth reported in this paper.

B ESTIMATING TRAINING-RUN VARIANCE

The bootstrap intervals in Table 4 resample evaluation documents and therefore capture evaluation sampling error alone. Variability attributable to training was estimated separately in an auxiliary experiment on the same evaluation suite.

Five instruction-tuning conditions, differing only in which synthetic mixture was blended into the instruction data, were each trained twice from the same base checkpoint under different seeds, and all ten resulting models were evaluated with the identical protocol of Appendix K. For each condition and each task, the absolute difference between the two seeds gives one observation of run-to-run variability. Over the fifty resulting condition and task pairs the mean absolute difference is 0.023, the median is 0.011 and the maximum is 0.135. Treating the two runs of a condition as independent draws implies a per-run standard deviation near 0.019, and therefore a standard deviation near 0.027 for a difference between two independently seeded conditions.

Combining that term with the evaluation term gives a standard deviation near 0.033 for the contrast between the two trace conditions in Section 7.4. The largest seed differences fall on the same task family that carries the trace-language contrast, which is why the contrast is reported in Section 7.5 as consistent with the controlled study and not as independently established.

C TASK DETAILS

BRANCHPROOF samples a target proof path of controlled depth and adds distractor branches that are locally coherent but do not support the queried answer. Shortcut variants add training-only marker correlations that identify the active branch with a chosen probability. Evaluation removes those markers.

ATTRCON samples objects, slots, values, and propagation constraints. The formal trace names slot/value relations directly. The natural trace describes the same propagation in controlled prose. This changes the symbolic surface while preserving the paired-rendering principle.

D TRAINING AND EVALUATION DETAILS

All main SFT runs use materialized datasets so that the formal and natural conditions see the same prompts and answers. Each train-depth subset contains 50k examples, and the online validation subset contains 256 examples. Main runs train for 10k optimizer steps with LoRA adapters on `q_proj`, `k_proj`, `v_proj`, `o_proj`, `up_proj`, `down_proj`, and `gate_proj`. The LoRA rank is 16, alpha is 32, dropout is 0.05, the optimizer learning rate is 10^{-4} , the per-device batch size is 1, gradient accumulation is 1 unless an ablation states otherwise, and training uses BF16 with gradient checkpointing.

Post-hoc evaluation uses vLLM with $\text{top-}p = 0.95$, temperature 1.0, 16 sampled generations per prompt, and $\text{pass@}k$ values $\{1, 2, 4, 8, 16\}$. The explicit depth grid is $\{1, 2, 5, 10, 12, 15, 18, 20, 25, 30, 35, 40, 45, 50\}$. For each seed and depth we evaluate 32 prompts, so a train-1-to-25 BRANCHPROOF long-depth band contains $5 \times 32 = 160$ prompts per seed. Reported uncertainty is seed-level standard deviation, not a confidence interval.

E TRACE GRAMMAR AND VALIDATION

Formal traces contain four blocks: constants, predicates, premises, and proof lines, followed by a conclusion and answer. A proof line contains a formula and a rule justification such as premise retrieval or implication elimination. The formal checker parses the generated blocks and verifies that each proof line is licensed by earlier lines and premises. Natural-language traces use controlled sentences for the same premises and proof steps. When possible, a task-specific translator maps those sentences back into the formal checker. Because translation coverage is task-specific, final-answer correctness is the primary cross-task metric and translated validity is reported as a diagnostic.

The operational BRANCHPROOF formal surface is:

```

702 <formal>
703 <constants> CONST_DEF+ </constants>
704 <predicates> PRED_DEF+ </predicates>
705 <premises> FORMULA+ </premises>
706 <proof> PROOF_LINE+ </proof>
707 <conclusion> ATOM </conclusion>
708 <answer> LABEL </answer>
709
710 CONST_DEF ::= CONST "=" LABEL
711 PRED_DEF  ::= PRED "x:" "x is" LABEL
712 ATOM      ::= PRED CONST
713 FORMULA   ::= ATOM | ATOM "&" ATOM "->" ATOM | ATOM "->" ATOM
714 PROOF_LINE ::= FORMULA "; R," REF
715             | ATOM "; ->E," REF "," REF
716             | ATOM "&" ATOM "; &I," REF "," REF

```

The ATTRCON surface uses the same block structure but replaces unary predicate atoms with object-slot-value relations. In both tasks, references must point to earlier proof lines or listed premises. Generated lines cannot justify future steps.

F GENERATOR AND EVALUATION MECHANICS

The BRANCHPROOF generator samples a target path, assigns labels to constants and predicates, then adds distractor branches that remain locally coherent but do not derive the queried conclusion. In pseudocode:

```

727 sample target answer label
728 sample initial object state
729 for depth step in 1..D:
730     sample active branch rule(s)
731     derive next object state
732     add coherent distractor rules for inactive branches
733 render one prompt
734 render the same latent proof as formal and natural-language targets

```

The ATTRCON generator follows the same paired-rendering principle with objects, slots, values, and propagation constraints. The natural-language target is not free-form prose. It is controlled text produced from the same latent derivation.

Answer extraction first searches for the final <answer> field. If no well-formed final answer is present, the generation is counted incorrect even when earlier text contains the gold label. Formal validation then parses block boundaries, declarations, formulas, proof-line references, and rule applications. Common validation failures are malformed block structure, undeclared symbols, references to unavailable lines, applying implication elimination to the wrong antecedent, and stopping before the queried conclusion. Natural-language validation adds one more possible failure mode: the controlled sentence must translate back into a formal step before the proof checker can evaluate it.

G PAIRED TARGET EXAMPLES

The following minimal examples show the paired target format used throughout the experiments. The prompt and answer are identical within each pair. Only the target trace substrate changes. Full experimental examples include more distractor premises and longer derivations.

BRANCHPROOF example 1.

```

754 Prompt:
755 a is ruby. a is north.

```

756 If a is north and a is ruby, then b is olive.
 757 What is b?
 758
 759
 760

761 Formal target

```
762 <formal>
763 <constants> a = a b = b </constants>
764 <predicates> Rx: x is ruby Nx: x is north
765 Ox: x is olive </predicates>
766 <premises> Ra Na Na & Ra -> Ob </premises>
767 <proof> Ra ; R,1
768 Na ; R,2
769 Na & Ra ; &I,5,4
770 Ob ; ->E,3,6 </proof>
771 <conclusion> Ob </conclusion>
772 <answer> olive </answer>
```

Natural-language target

```
<think>
a is ruby. a is north.
Since a is north and a is ruby, b is olive.
Therefore, b is olive.
</think>
<answer> olive </answer>
```

774 BRANCHPROOF example 2.

775
 776
 777 Prompt:
 778 a is cobalt. a is east.
 779 If a is east and a is cobalt, then b is maple.
 780 If b is maple, then b is south.
 781 If b is south and b is maple, then c is cedar.
 782 What is c?
 783
 784
 785

786 Formal target

```
787 <formal>
788 <constants> a = a b = b c = c </constants>
789 <predicates> Cx: x is cobalt Ex: x is east
790 Mx: x is maple Sx: x is south Dx: x is cedar </predicates>
791 <premises> Ca Ea
792 Ea & Ca -> Mb
793 Mb -> Sb
794 Sb & Mb -> Dc </premises>
795 <proof> Ca ; R,1
796 Ea ; R,2
797 Ea & Ca ; &I,7,6
798 Mb ; ->E,3,8
799 Sb ; ->E,4,9
800 Sb & Mb ; &I,10,9
801 Dc ; ->E,5,11 </proof>
802 <conclusion> Dc </conclusion>
803 <answer> cedar </answer>
```

Natural-language target

```
<think>
a is cobalt. a is east.
Since a is east and a is cobalt, b is maple.
Since b is maple, b is south.
Since b is south and b is maple, c is cedar.
Therefore, c is cedar.
</think>
<answer> cedar </answer>
```

803 ATTRCON example 1.

804
 805
 806
 807 Prompt:
 808 s0 has value red. s1 has value blue.
 809 If s0 is red and s1 is blue, then s2 is green.
 Which values fill all constrained slots?

Formal target

```
<formal>
<constants> s0 = attribute slot 0 s1 = attribute
s2 = attribute slot 2
red = attribute value red
blue = attribute value blue
green = attribute value green </constants>
<predicates> Value(x,y): slot x has value y
Constraint(x,y,z,w,u,v): values y and w at slots
jointly force value v at slot u </predicates>
<premises> Value(s0,red) Value(s1,blue)
Constraint(s0,red,s1,blue,s2,green)
Value(s0,red) & Value(s1,blue) &
Constraint(s0,red,s1,blue,s2,green)
-> Value(s2,green) </premises>
<proof> Value(s0,red) ; R,1
Value(s1,blue) ; R,2
Constraint(s0,red,s1,blue,s2,green) ; R,3
Value(s0,red) & Value(s1,blue) ; &I,5,6
Value(s0,red) & Value(s1,blue) &
Constraint(s0,red,s1,blue,s2,green) ; &I,8,7
Value(s2,green) ; ->E,4,9 </proof>
<answer> red-blue-green </answer>
```

Natural-language target

```
<think>
s0 has value red.
s1 has value blue.
The applicable joint constraint maps s0=red and s1=blue
to s2=green.
Combine the two prerequisite slot values.
Therefore s2 has value green.
x/answer>
<answer> red-blue-green </answer>
```

ATTRCON example 2.

Prompt:

s0 has value amber. s1 has value ivory.
If s0 is amber and s1 is ivory, then s2 is olive.
If s0 is amber and s2 is olive, then s3 is coral.
Which values fill all constrained slots?

Formal target

```
<formal>
<constants> s0 = attribute slot 0 s1 = attribute
s2 = attribute slot 2 s3 = attribute slot 3
amber = attribute value amber
ivory = attribute value ivory
olive = attribute value olive
coral = attribute value coral </constants>
<predicates> Value(x,y): slot x has value y
Constraint(x,y,z,w,u,v): values y and w at slots
jointly force value v at slot u </predicates>
<premises> Value(s0,amber) Value(s1,ivory)
Constraint(s0,amber,s1,ivory,s2,olive)
Value(s0,amber) & Value(s1,ivory) &
Constraint(s0,amber,s1,ivory,s2,olive) -> Value(s2,olive)
Constraint(s0,amber,s2,olive,s3,coral)
Value(s0,amber) & Value(s2,olive) &
Constraint(s0,amber,s2,olive,s3,coral)
-> Value(s3,coral) </premises>
<proof> Value(s0,amber) ; R,1
Value(s1,ivory) ; R,2
Constraint(s0,amber,s1,ivory,s2,olive) ; R,3
Value(s0,amber) & Value(s1,ivory) ; &I,7,8
Value(s0,amber) & Value(s1,ivory) &
Constraint(s0,amber,s1,ivory,s2,olive) ; &I,10,9
Value(s2,olive) ; ->E,4,11
Constraint(s0,amber,s2,olive,s3,coral) ; R,5
Value(s0,amber) & Value(s2,olive) ; &I,7,12
Value(s0,amber) & Value(s2,olive) &
Constraint(s0,amber,s2,olive,s3,coral) ; &I,14,13
Value(s3,coral) ; ->E,6,15 </proof>
<answer> amber-ivory-olive-coral </answer>
```

Natural-language target

```
<think>
s0 has value amber.
s1 has value ivory.
The applicable joint constraint maps s0=amber and s1=ivory
to s2=olive.
Combine the two prerequisite slot values.
Therefore s2 has value olive.
The applicable joint constraint maps s0=amber and s2=olive
to s3=coral.
Combine the two prerequisite slot values.
Therefore s3 has value coral.
</think>
<answer> amber-ivory-olive-coral </answer>
```

H METRIC DEFINITIONS

Answer correctness is computed after extracting the final answer field from the generation. Formal joint validity requires both answer correctness and a mechanically valid formal trace. Natural-language joint validity requires answer correctness and successful translation of the generated natural-language trace into a valid formal proof. Because the natural-language translator is task-specific, answer correctness is the primary cross-task metric.

Table 5: Rendered document lengths under the Qwen2.5-7B tokenizer, 72,000 documents per corpus, and packing statistics at sequence length 8,192.

corpus	p50	mean	p99	max	windows	real tokens	pack. eff.
Logic	3,749	3,856	7,673	7,727	35,492	277,732,923	0.9551
Natural	3,819	3,907	7,744	7,851	35,979	281,372,725	0.9545
Condensed	1,450	1,500	3,007	3,014	13,311	108,099,330	0.9912
LongDoc	3,712	3,832	n/a	7,750	n/a	n/a	n/a

I LENGTH AND TOKEN-BUDGET CONTROLS

Formal and natural renderings of the same latent proof differ in length, which raises the possibility that the controlled comparison measures the quantity of supervision. Three controls bound that explanation.

Generation budgets are identical by construction. Both conditions receive a 7,168-token generation limit inside a 16,384-token window. A tokenizer audit over the full evaluation set confirms that every gold prompt and every gold target fits inside that limit in both renderings, so neither condition is truncated at generation time and neither is advantaged by a looser budget.

Supervised token counts are equated in a matched-token training run, in which the number of target tokens per condition is held equal by subsampling the longer condition. The ordering of the two conditions is unchanged.

The transfer study of Section 7 supplies a third control at a different scale. The **LongDoc** condition matches the token-length histogram of the formal traces to within one percent at the median while carrying no deep deduction, and it does not reproduce the gain of the trace conditions on the external benchmark. Length alone therefore does not account for the effect in either setting.

J BOUNDARY CONDITION DETAILS

Capacity. Single-modality training at three scales gives out-of-distribution answer pass@1 of 69.9 ± 11.1 formal against 74.9 ± 7.0 natural on Qwen2.5-7B, and 63.9 ± 5.1 formal against 98.9 ± 0.8 natural on OLMo-3-32B. Formal answer pass@16 at 32B is 98.8 ± 0.9 , so the deficit at 32B is a deficit of probability mass and not of coverage.

Mode conditioning. Concatenating both renderings into a single target yields 0.5 ± 0.4 percent out-of-distribution pass@1 with the formal rendering first and 0.2 ± 0.1 percent with the natural rendering first. Mode-conditioned training with an output-mode tag reaches 85.0 ± 8.7 answer pass@1 in formal mode at 32B and 77.2 ± 11.7 in natural mode. Raw generations show clean separation between the modes.

Shortcuts. A spurious schema marker planted on 80 percent of training examples and removed at test time gives 43.1 ± 7.7 out-of-distribution pass@1 for formal against 78.5 ± 14.9 for natural. Additional marker types and injection rates were measured and are not reported here.

Scale floor. From-scratch training at 50M to 200M parameters on the same corpora gives out-of-distribution joint pass@k of 0.0 for all $k \leq 8$ at every size in both conditions.

K MIDTRAINING CONFIGURATION

Optimisation. Every condition trains Qwen2.5-7B for 2,385 steps at sequence length 8,192 with a global batch of 128 sequences, formed as micro-batch 2 with gradient accumulation 32 over eight A100-80 devices. The resulting budget is 2,500,853,760 tokens per condition, chosen to equal $4,770 \times 128 \times 4,096$ so that it matches the token budget of an earlier study at half the sequence length. Learning rate is 1×10^{-5} with minimum 1×10^{-6} , 128 warmup steps, decay beginning at step 128 and a decay horizon of 18,946 steps. Checkpoints are written every 250 steps. Measured throughput

is approximately 36.0 seconds per iteration and 29,000 tokens per second, at 191 TFLOP/s per device. Collective-operation timeouts are set to 120 minutes. A run exceeds the 24-hour scheduling limit and is completed as two serialized allocations, with the second resuming from the newest complete checkpoint.

Replaced slice. The 72,000 source examples are generated at seed 20260907 for the reinforcement-learning study and at seed 20260830 for the midtraining study, with depths 1 to 25 sampled round-robin, branching factor 4, distractor ratio 0.5 and the `hard_fsa_schema` variant. Seeds are disjoint from those used for evaluation corpora and for sampling. A per-condition document weight is solved so that the realised synthetic share of loss tokens is exactly ten percent after padding is accounted for, which is necessary because the three renderings have different packing efficiencies.

Packing. Documents are packed by a document-preserving packer at sequence length 8,192 under a fixed shuffle seed of 42. The packer never splits a document. A document exceeding one window is excluded and counted as overlength. Tail padding uses a dedicated padding token that is masked out of the loss. Overlength and split counts are zero for every condition, reported in Table 5. Before training begins, a decoded-batch audit reconstructs real loader windows and verifies that no split document is present, that the padding token is masked, and that the realised replacement share matches the target. Training refuses to start if any gate fails.

Instruction tuning. Each midtrained checkpoint is converted to the Hugging Face format and instruction-tuned on 100,000 examples of a public instruction corpus at a pinned revision, for one epoch at learning rate 5×10^{-6} with warmup ratio 0.03, maximum sequence length 4,096, per-device batch size 1, and seed 3407. Training is full-parameter under fully sharded data parallelism with gradient checkpointing. Only the final checkpoint is written.

Evaluation. Greedy evaluation uses a standard harness with a paged-attention inference backend in `bfloat16`, a maximum model length of 8,192, a chat template applied to every prompt, and per-sample logging enabled. The reported metric is exact match on the extracted answer. The external family is ProofWriter under the open-world assumption at inference depths 0, 1, 2, 3 and 5 with 500 items per depth. The near-transfer family is a held-out sample from the same generator at depths 5, 10, 15, 20 and 25 with 200 items per depth.

Sampled evaluation draws sixteen completions per item at temperature 0.8 and top- p 0.95 under a fixed seed, with a 2,048-token response limit. Prompts are replayed verbatim from the logged greedy samples, so the greedy and sampled measurements differ in the sampler alone and in nothing else. The scorer used for sampled generations imports the same extraction and matching function used by the greedy harness, and was verified to reproduce all fifteen published greedy exact-match values before any sampled number was computed. Audits require the expected item count per task, no empty generations beyond a tolerance, and no truncated prompts.

Statistics. Contrasts in Table 4 come from a paired bootstrap over evaluation documents with 4,000 resamples, resampling document indices once per replicate and applying the same indices to every condition and depth. Reported intervals are the 2.5th and 97.5th percentiles of the replicate distribution. Training-run variance is estimated separately from an earlier campaign on the same evaluation suite in which five conditions were each trained under two instruction-tuning seeds, giving fifty condition and task pairs.

L PROOF OF PROPOSITION 1

Proof. The samples are independent, so the index of the first accepted sample is independent of whether that sample is correct. Acceptance occurs at least once with probability $1 - (1 - a - b)^k$. Conditioned on acceptance, the accepted sample is drawn from the acceptance event, in which correctness has probability $a/(a + b)$. Substituting $b = 0$ gives $1 - (1 - a)^k$. $\square$

M EXTENDED TESTBED DESCRIPTION

Each example begins with a latent directed acyclic proof. The generator emits a prompt containing ground facts, rules, a query and distractor branches, then traverses the supporting path to produce the target answer. A rendering function writes that path in one of two trace languages. The `FORMAL` rendering declares constants and predicates and applies named inference rules to numbered lines. The `NATURAL` rendering expresses the same substitutions and conclusions, line for line, in controlled sentences. Both targets end with the same answer field, and the prompt is natural language in both conditions.

`BRANCHPROOF` tests composition across a chain of Horn-style inferences in which every step offers four locally plausible branches, exactly one of which is derivable from the current state. Every layer introduces a fresh constant. The generator computes the full Horn closure of each example and accepts it only when exactly one candidate answer is derivable, which removes the possibility that a model is rewarded for guessing among several defensible answers. `ATTRCON` propagates values among object attributes under generated constraints. Its vocabulary and proof construction are disjoint from `BRANCHPROOF`, so it provides an independent test of the same hypothesis. Appendix G gives complete paired targets, Appendices E and F specify the grammar, the generator and the validators, and Appendix A documents the answer-uniqueness audit of the final generator.

Generation parameters are held fixed across every controlled experiment reported below. The branching factor is 4, the distractor ratio is 0.5, and the schema variant is the one denoted `hard_fsa_schema` in the release. Depths are sampled round-robin so that each depth in the training range contributes an equal number of examples. Seeds are recorded per corpus and are disjoint across roles, so that no corpus used for training shares a seed with a corpus used for evaluation or for sampling.